



University of
Pittsburgh

Algorithms and Data Structures 2

CS 1501



Fall 2022

Sherif Khattab

ksm73@pitt.edu

(Slides are adapted from Dr. Ramirez's and Dr. Farnan's CS1501 slides.)

Announcements

- Upcoming Deadlines
 - Homework 3: this Friday @ 11:59 pm
 - Lab 2: next Monday @ 11:59 pm
 - Assignment 1: Monday Oct 10th @ 11:59 pm
- Please include all instructors when sending private messages on Piazza, if possible
- **Student Support Hours** of the teaching team are posted on the Syllabus page

Previous lecture

- Red-Black BST (self-balancing BST)
 - definition and basic operations
 - delete
 - runtime of operations
- Turning recursive tree traversals to iterative

Muddiest Points

- **Q: So, just to recap: When we add a new node to a Red-Black BST, if that leads to a violation we need to fix that violation and all others back up to the tree's root before we can say we're done?**
- **Yep!**

Muddiest Points

- **Q: I think there may be an overlap in which sections muddiest points you are displaying**
- Yes. But I made sure to explain the concept before going over the muddies points


Muddiest Points

- **Q: how does a tree map work? a hashmap works effectively by 'converting' a string to an int by using a hash code to map a key to a value, what is a treemap and how does it do that?**
- Both TreeMap and HashMap implement the Map or Dictionary interface
- Nothing in the Map interface requires the conversion into an integer
- TreeMap uses a Binary Search Tree instead of a hash table to perform add, search, delete, etc.

Muddiest Points


- **Q: Could you please explain the iterative approach to the BST is again?**
- **Sure.**

BST search using iteration



→ $x = \text{root}$
 $\text{while}(x \neq \text{null})$
 if equal break;
 if $<$ $x = x.\text{left}$
 if $>$ $x = x.\text{right}$
}


BST search using iteration



$x = \text{root}$
→ while($x \neq \text{null}$) {
 if equal break;
 if $<$ $x = x.\text{left}$
 if $>$ $x = x.\text{right}$
}

Search for 19


BST search using iteration



```
x = root
while(x != null){
    if equal break;
    if < x = x.left
    if > x = x.right
}
```

Search for 19

BST search using iteration




x



```
x = root
while(x != null){
    if equal break;
    if < x = x.left
    if > x = x.right
}
```

Search for 19


BST search using iteration



```
x = root
while(x != null){
    if equal break;
    if < x = x.left
    if > x = x.right
}
```

Search for 19

BST search using iteration



```
x = root  
while(x != null){  
    → if equal break;  
    if < x = x.left  
    if > x = x.right  
}
```

Search for 19

Muddiest Points

- **Q: The iterative method for pre-order was confusing for me to grasp but it started to make more sense once we did in-order and post-order.**
- **Thanks!**

Muddiest Points

- **Q: How does a Red Black BST keep track of the data in the nodes if it's just comparing colors?**
- It is comparing data as we go down the tree (except for delete) and checking colors as we climb back up the tree

Muddiest Points

- **Q: can you explain the stack and what pushing/popping are?**
- Stack is an abstract data type in which data items are ordered in a Last In First Out order.
- Pushing into a Stack means to add an item at the “top” of the stack
- Popping means to remove the top item

This Lecture

- Binary Search Tree uses comparisons between keys to guide the searching
- What if we use the digital representation of keys for searching instead?
 - Keys are represented as a sequence of digits (e.g., bits) or alphabetic characters
- Digital Searching Problem

Digital Searching Problem

- Input:
 - a (large) dynamic set of data items in the form of
 - n (key, value) pairs; key is a string from an alphabet of size R
 - Each key has b bits or w characters (the chars are from the alphabet)
 - What is the relationship between b and w ?
 - a *target key* (k)
- Output:
 - The corresponding value to k if target key found
 - Key not found otherwise

Digital Search Trees (DSTs)

Instead of looking at less than/greater than, let's go left or right based on the bits of the key

So, we again have 4 options:

- current node is null, k not found
- k is equal to the current node's key, k is found, return corresponding value
- current bit of k is 0, continue to left child
- current bit of k is 1, continue to right child

DST example: Insert and Search

Insert:

4 0100

3 0011

2 0010


6 0110

5 0101

Search:


3 0011

7 0111




DST and Prefixes

- In a DST, each node shares a **common prefix** with all nodes in its subtree
 - E.g., 6 shares the prefix "01" with 5
- In-order traversal doesn't produce a sorted order of the items
 - Insertion algorithm can be modified to make a DST a BST at the same time




DST example: Delete

- Delete 3
- Can replace it with any leaf in its subtree
- Let's replace it with 2
- OK because 2 shares "0" as a prefix with 3, so it also shares "0" as a prefix with 6 and 5



DST example: Delete

- Delete 3
- Can replace it with any leaf in its subtree
- Let's replace it with 2
- OK because 2 shares "0" as a prefix with 3, so it also shares "0" as a prefix with 6 and 5




DST example: Variable length keys

- Insert

1 01

- Must be in place of 6
- Replace 6 by 1 and re-insert 6



DST example: Variable length keys


- Insert

1 01

- Must be in place of 6

- Replace 6 by 1 and re-insert

6 0110



Analysis of digital search trees

$$\text{average Case runtime} = \sum_{\text{all cases}} P_r(\text{Case}_i) \times \text{runtime for Case}_i$$

- Runtime?
 - $O(b)$, b is the bit length of the target or inserted key
 - On average, $b = \log(n)$
 - When branching according to a 0 or 1 is equally likely
 - In general $b \geq \lceil \log n \rceil$
- We end up doing many **equality** comparisons against the full key
- This is better than less than/greater than comparison in BST
- Can we improve on this?

Radix search tries (RSTs)

- Trie as in re**trie**ve, pronounced the same as “try”
- Instead of storing keys inside nodes in the tree, we store them implicitly as paths down the tree
 - Interior nodes of the tree only serve to direct us according to the bitstring of the key
 - Values can then be stored at the end of key’s bitstring path (i.e., at leaves)
 - RST uses less space than BST and DST

```

graph TD
    Root(( )) ---|0| Node0(( ))
    Root ---|1| Node1(( ))
    Node0 ---|0| V2((V2))
    Node0 ---|1| Node0_1(( ))
    Node0_1 ---|0| V3((V3))
    Node0_1 ---|1| V4((V4))
    Node1 ---|0| Node1_0(( ))
    Node1 ---|1| Node1_1(( ))
    Node1_0 ---|0| V5((V5))
    Node1_0 ---|1| V6((V6))
    Node1_1 ---|0| V7((V7))
    Node1_1 ---|1| V8((V8))
  
```

4 0100

3 0011

2 0010


6 0110

5 0101

3 0011

7 0111

Binary RST




RST analysis

- Runtime?
- $O(b)$, the bit length of the key
 - However, this time we don't have full key comparisons
- Would this structure work as well for other key data types?
- Characters?
 - Characters are the same as 8-bit ints (assuming simple ascii)
- Strings?
- May have huge bit lengths
- How to store Strings?

Larger branching factor tries

- In our binary-based Radix search trie, we considered one bit at a time
- What if we applied the same method to characters instead of bits in a string?
 - What would this new structure look like?
 - How many children per node?
 - up to R (the alphabet size)
 - Also called R -way radix search tries
- Let's try inserting the following strings into an trie:
 - she, sells, sea, shells, by, the, sea, shore

Another trie example



Analysis

- Runtime?
- $\Theta(w)$ where w is the character length of the string
 - So what do we really gain over RSTs?
 - For strings, $w < b$, and overall tree height is reduced
 - $W = \frac{b}{\log R}$ where R is the alphabet size
 - For binary RST: average tree height = $\log_2(n)$
 - For R -way RST: average tree height = $\log_R(n)$

Further analysis


- Search Miss
 - Require an average of $\log_R(n)$ nodes to be examined
 - Where R is the size of the alphabet being considered
 - Proof in Proposition H of Section 5.2 of the text
 - Average # of checks with 2^{20} keys in an RST?
 - $\log_2 n = \log_2 2^{20} = 20$
 - With 2^{20} keys in a large branching factor trie, assuming 8-bits at a time?
 - $\log_R n = \log_{256} 2^{20} = \log_{256} (2^8)^{2.5} = \log_{256} 256^{2.5} = 2.5$

Implementation Concerns

- See TrieSt.java
 - Implements an R-way trie
- Basic node object:


Where R is the branching factor

```
private static class Node {  
    private Object val;  
    private Node[] next = new Node[R];  
}
```




- Non-null **val** means we have traversed to a valid key
- Again, note that keys are not directly stored in the trie at all

R-way trie example



R-way RST




Summary of running time

	insert	Search hit	Search miss
binary RST	$\Theta(b)$	$\Theta(b)$	$\Theta(\log_2 n)$ on average
multi-way RST	$\Theta(w)$	$\Theta(w)$	$\Theta(\log_R n)$

So what's the catch with R-way RST?


- Space!
 - Considering 8-bit ASCII, each node contains 2^8 references!
 - This is especially problematic as in many cases, a lot of this space is wasted
 - Common paths or prefixes for example, e.g., if all keys begin with "key", that's 255×3 wasted references!
 - At the lower levels of the trie, most keys have probably been separated out and reference lists will be sparse



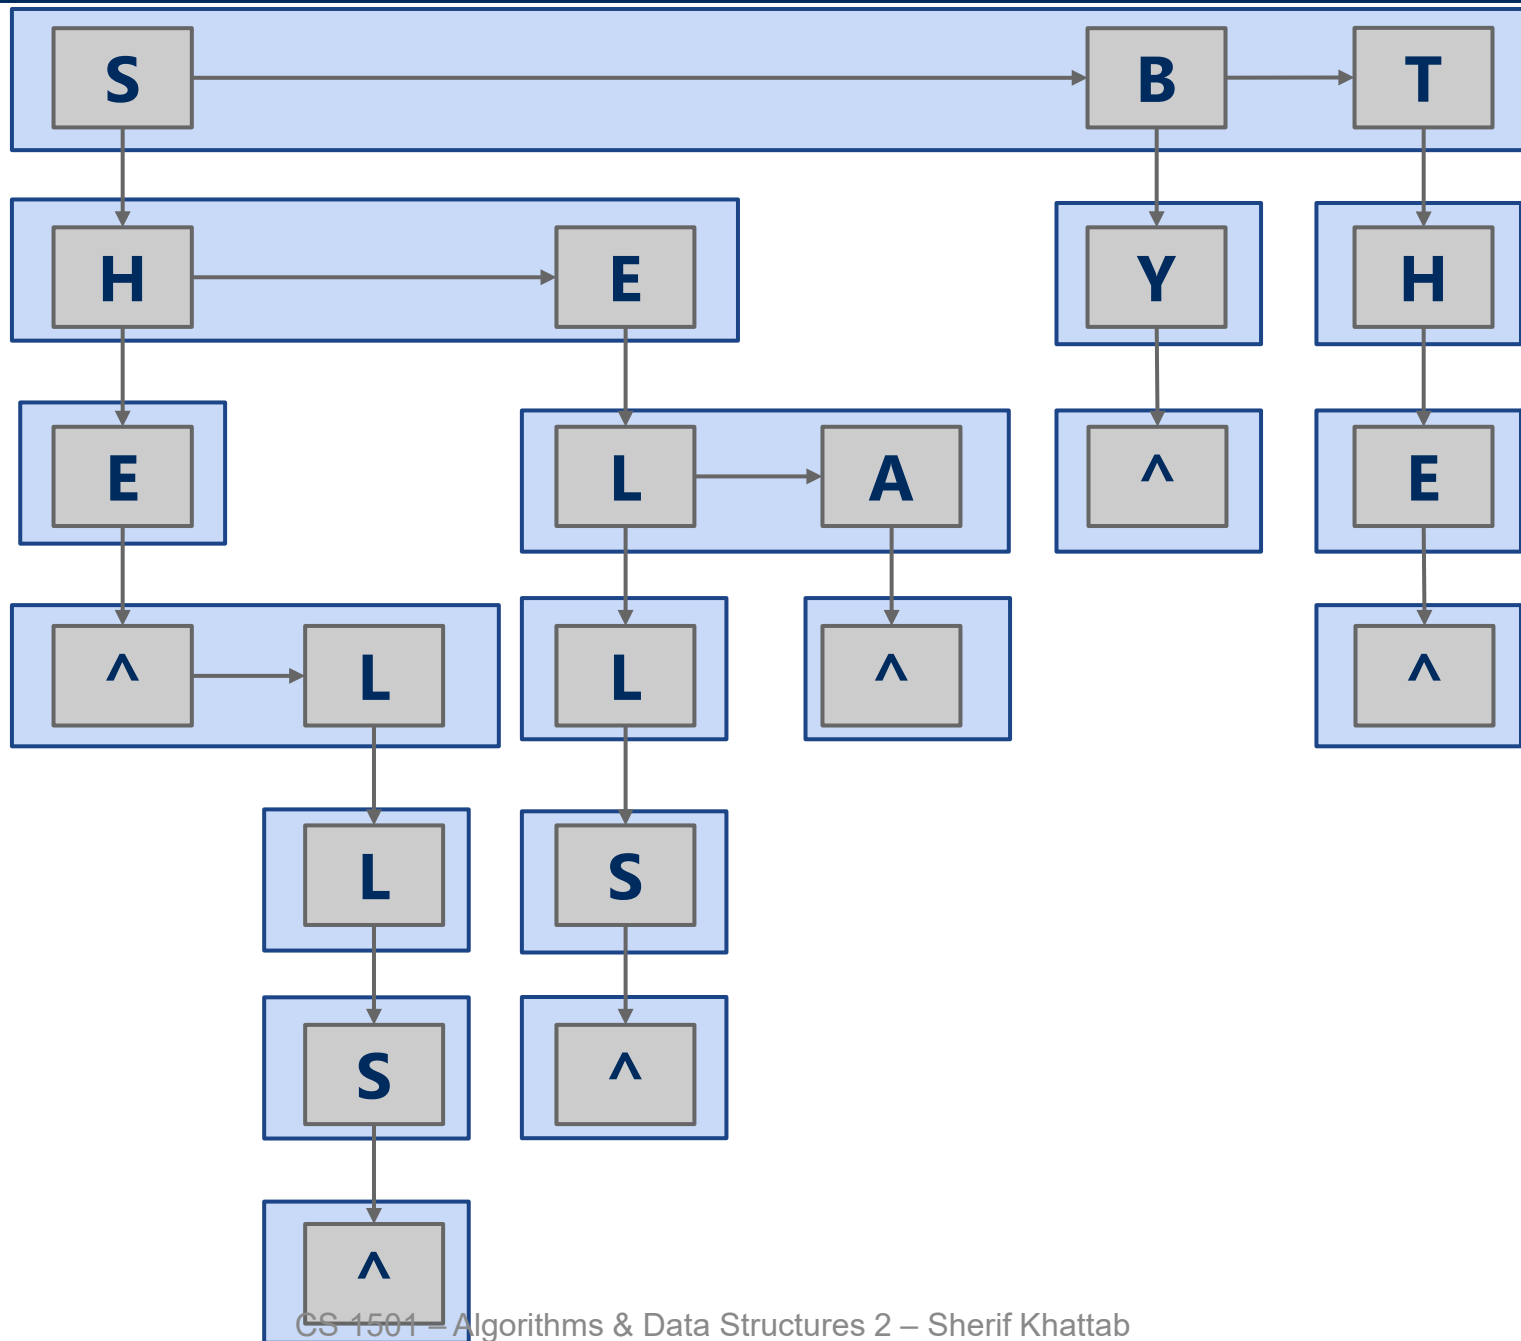
De La Briandais tries (DLBs)

Main idea: replace the .next array of the R-way trie with a linked-list

DLB Example



DLB Example 2: nodes vs. nodelets



DLB analysis


- How does DLB performance differ from R-way tries?
- Which should you use?

		Search hit insert	
R-way RST		$\theta(w)$	
	DLB	$\theta(wR)$	


R-way RST vs. DLB

aaa
bbb
ccc

R-way
RST



DLB



$$\theta(w, 1)$$

characters in the key

$$\Theta(W \cdot R)$$

Runtime Comparison for Search Trees/Tries

	Search hit	Search miss <i>(average)</i>	insert
BST	$\Theta(n)$	$\Theta(\log n)$	$\Theta(n)$
RB-BST	$\Theta(\log n)$	$\Theta(\log n)$	$\Theta(\log n)$
DST	$\Theta(b)$	$\Theta(\log n)$	$\Theta(b)$
RST	$\Theta(b)$	$\Theta(\log n)$	$\Theta(b)$
R -way RST	$\Theta(w)$	$\Theta(\log n)$	$\Theta(w)$
DLB	$\Theta(w \cdot R)$	$\Theta(\log_{\substack{R \\ n}} n \cdot R)$	$\Theta(w \cdot R)$

Final notes on Search Tree/Tries

- We did not present an exhaustive look at search trees/tries, just the sampling that we're going to focus on
- Many variations on these techniques exist and perform quite well in different circumstances
 - Ternary search Tries
 - R-way tries without 1-way branching
- See the table at the end of Section 5.2 of the text