



University of  
Pittsburgh

# Algorithms and Data Structures 2

## CS 1501



Spring 2023

Sherif Khattab

ksm73@pitt.edu

(Slides are adapted from Dr. Ramirez's and Dr. Farnan's CS1501 slides.)

# Announcements

- Upcoming Deadlines
  - Homework 11: this Friday @ 11:59 pm
  - Lab 9: Tuesday 4/4 @ 11:59 pm
  - Assignment 4: Friday 4/14 @ 11:59 pm
    - Support video and slides on Canvas + Solution for Labs 8 and 9

# Previous lecture

- Priority Queue ADT
  - Heap implementation

# This Lecture

- Minimum Spanning Tree (MST)
  - Prim's MST algorithm
    - naïve implementation
    - Best Edges array implementation
    - using a min-heap
  - Kruskal's MST algorithm
- Weighted Shortest Paths problem
  - Dijkstra's single-source shortest paths algorithm
  - Bellman-Ford's shortest paths algorithm

# Neighborhood connectivity Problem

- We want to keep a set of neighborhoods **connected** with the **minimum cost** possible
- **Input:** A set of neighborhoods and a file:
  - neighborhood i, neighborhood j, cost of connecting the two neighborhoods
  - ...
- **Output:** A set of neighborhood pairs to be connected and a total cost such that
  - We can go from any neighborhood to any other (**connected**)
  - The total cost should be minimum (i.e., as small as it can be) (**minimal cost**)

# Think Data Structures First!

- How can we structure the input in computer memory?
- Can we use Graphs?
- What about the costs? How can we model that?

## We said spatial layouts of graphs were irrelevant

- We define graphs as sets of vertices and edges
- However, we'll certainly want to be able to reason about **bandwidth, distance, capacity**, etc. of the real world things our graph represents
  - Whether a link is 1 gigabit or 10 megabit will drastically affect our analysis of traffic flowing through a network
  - Whether a road is single-lane or 4-lane
  - Whether two airports are 2000 miles apart or 200 miles apart will have an effect on the number of flights going in and out between them

# We can represent such information with edge weights

- How do we store edge weights?
  - Adjacency matrix?
  - Adjacency list?
  - Do we need a whole new graph representation?



# New problems!

How do weights affect finding spanning trees and shortest paths?

- The weighted variants of these problems are called ***minimum spanning tree*** and ***weighted shortest path***


# Minimum spanning trees (MST)

- Graphs can potentially have multiple spanning trees
  - BFS, DFS traversals find possible different spanning trees
- MST is the spanning tree that has the **minimum sum** of the weights of its edges

# Prim's algorithm

- Initialize  $T$  to contain the starting vertex
  - $T$  will eventually become the MST
- While there are vertices not in  $T$ :
  - Find a **minimum edge-weight edge** that connects a vertex **in  $T$**  to a vertex **not yet in  $T$**
  - Add the edge with its vertex to  $T$

# Prim's algorithm



# Runtime of Prim's

- At each step, check all possible edges
- For a complete graph:
  - First iteration:
    - $v - 1$  possible edges
  - Next iteration:
    - $2(v - 2)$  possibilities
      - Each vertex in  $T$  shared  $v-1$  edges with other vertices, but the edges they shared with each other already in  $T$
  - Next:
    - $3(v - 3)$  possibilities
  - ...
- Runtime:
  - $\sum_{i=1}^{v-1} (i * (v - i)) = \Theta(\text{largest term} * \text{number of terms})$
  - number of terms =  $v-1$
  - largest term is  $v^2/4$  (when  $i=v/2$ )
  - Evaluates to  $\Theta(v^3)$


# Do we need to look through all remaining edges?

- No! We only need to consider the ***best edge*** possible for each vertex!
- The best edge of a vertex is the edge with the **minimum weight** connecting to the vertex from a vertex already in T
  - Adding a vertex to T → new edges become
  - Best edge values can be **updated** as we add vertices to T

# An enhanced implementation of Prim's Algorithm

- Add start vertex to T
- Repeat until all vertices added to T
  - Check the **neighbors of the added** vertex
    - update best edge values if needed
    - update **parent** as well
  - Add to T a vertex with the **smallest best edge**

# Prim's algorithm



	0	1	2	3	4	5	6
Parent:	--	0	0	2	5	2	5
Best Edge:	0	4	8	2	5	1	2



# Runtime of the Best Edges Implementation

- For every vertex we add to  $T$ , check and possibly update neighbors
- Let's assume we use an **adjacency matrix**:
  - Takes  $\Theta(v)$  to check the neighbors of a given vertex
  - Time to update parent/best edge arrays?
    - $\Theta(1)$
  - Time to pick next vertex?
    - $\Theta(v)$
- Total:  $v * 2 \Theta(v) = \Theta(v^2)$

# Runtime of the Best Edges Implementation

- For every vertex we add to  $T$ , check and possibly update neighbors
- Let's assume we use an **adjacency lists**:
  - Takes  $\Theta(d)$  to check the neighbors of a given vertex
  - Time to update parent/best edge arrays?
    - $\Theta(1)$
  - Time to pick next vertex?
    - $\Theta(v)$
- Total:  $v * 2 \Theta(v) = \Theta(v^2)$

# Prim's MST Algorithm

- seen, parent, and BestEdge are arrays of size  $v$
- Initialize seen to false, parent to -1, and BestEdge to infinity
- BestEdge[start] = 0
- for  $i = 0$  to  $v-1$ 
  - Find  $w$  s.t. seen[ $w$ ] = false and BestEdge[ $w$ ] is minimum over all unseen vertices
  - seen[ $w$ ] = 1
  - for each neighbor  $x$  of  $w$ 
    - if(BestEdge[ $x$ ] > edge weight of edge ( $w, x$ )
      - BestEdge[ $x$ ] = edge weight of ( $w, x$ )
      - parent[ $x$ ] =  $w$
- The parent array represents the found MST

# Prim's MST Algorithm

- seen, parent, and BestEdge are arrays of size  $v$
  - Initialize seen to false, parent to -1, and BestEdge to infinity
  - BestEdge[start] = 0
  - for  $i = 0$  to  $v-1$ 
    - **Find  $w$  s.t. seen[ $w$ ] = false and BestEdge[ $w$ ] is minimum over all unseen vertices**
    - seen[ $w$ ] = 1
    - for each neighbor  $x$  of  $w$ 
      - if(BestEdge[ $x$ ] > edge weight of edge ( $w$ ,  $x$ )
        - BestEdge[ $x$ ] = edge weight of ( $w$ ,  $x$ )
        - parent[ $x$ ] =  $w$
- The parent array represents the found MST

# What about a faster way to pick the best edge?

- Sounds like a job for a priority queue!
  - Priority queues can remove the min value stored in them in  $\Theta(\log n)$ 
    - Also  $\Theta(\log n)$  to add to the priority queue

# Let's maintain best edge values in a PQ!

- PQ will need to be **indexable** to **update** the best edge
- This is the idea of *eager Prim's*

# Eager Prim's Runtime


- $v$  inserts
  - $v \log v$
- $e$  updates
  - $e \log v$
- $v$  removeMin
  - $v \log v$
- Total:  $(e+v) \log v$
- Assuming connected graph
  - $e \geq v - 1$
- $e+v = \Theta(e)$
- Total runtime =  $e \log v$

## We can be a bit lazy: let's keep $e$ edges in the PQ

- PQ doesn't have to be indexable
- Lazy Prim's implementation
  - Visit a vertex
  - Add edges coming out of it to a PQ
  - While there are unvisited vertices, pop from the PQ for the next vertex to visit and repeat



# Prim's with a priority queue



PQ:

1: (0, 2)

2: (5, 3)

3: (1, 4)

4: (2, 5)

5: (2, 3)

5: (0, 3)

5: (2, 1)

6: (0, 1)

6: (2, 4)

6: (5, 4)

# Runtime using a priority queue

- Have to insert all  $e$  edges into the priority queue
  - In the worst case, we'll also have to remove all  $e$  edges
- So we have:
  - $e * \Theta(\lg e) + e * \Theta(\lg e)$
  - $= \Theta(2 * e \lg e)$
  - $= \Theta(e \lg e)$
- This algorithm is known as *lazy Prim's*

# Comparison of Prim's implementations

- Parent/Best Edge array Prim's

- Runtime:  $\Theta(v^2)$
- Space:  $\Theta(v)$

- Lazy Prim's

- Runtime:  $\Theta(e \lg e)$
- Space:  $\Theta(e)$
- Requires a PQ

- Eager Prim's

- Runtime:  $\Theta(e \lg v)$
- Space:  $\Theta(v)$
- Requires an indexable PQ


How do these compare?



# Another MST algorithm

- Kruskal's MST:
  - Insert all edges into a PQ
  - Grab the min edge from the PQ that does not create a cycle in the MST
  - Remove it from the PQ and add it to the MST

# Kruskal's example



PQ:

1: (0, 2)

2: (3, 5)

3: (1, 4)

4: (2, 5)

5: (2, 3)

5: (0, 3)

5: (1, 2)

6: (0, 1)

6: (2, 4)

6: (4, 5)

# Kruskal's runtime

- Instead of building up the MST starting from a single vertex, we build it up using edges all over the graph
- How do we efficiently implement cycle detection?
  - BFS/DFS
    - $v + e$
  - Union/Find data structure
    - $\log v$

# Kruskal's Runtime

- $e$  iterations
  - removeMin
    - $\log e$
  - Cycle detection
    - $v + e$  using DFS/BFS
    - $\log v$  using Union/Find
- Total:  $e \log e$
- Assuming connected graph
  - $v - 1 \leq e \leq v^2$
  - $\log v \leq \log e \leq 2 \log v$
  - $\log e = \Theta(\log v)$
- Total runtime:  $e \log v$
- Same as Prim's

# Problem of the Day: Weighted Shortest Paths


- Input:
  - A road network
    - Road segments and intersections
    - Road segments are labeled by travel time
      - From length and maximum speed
      - How do we get max speed?
  - Starting address and destination address
- Output:
  - A shortest path from source to destination



# Dijkstra's algorithm

- Set a distance value of `Double.POSITIVE_INFINITY` for all vertices
- $\text{distance}[\text{start}] = 0$
- Set  $\text{cur} = \text{start}$
- While destination is not visited:
  - For each unvisited neighbor  $x$  of  $\text{cur}$ :
    - Compute distance from start to  $x$  through  $\text{cur}$ 
      - $\text{distance}[\text{cur}] + \text{weight of edge between cur and } x$
    - Update  $\text{distance}[x]$  if computed distance  $< \text{distance}[x]$
  - Mark  $\text{cur}$  as visited
  - Let  $\text{cur}$  be the unvisited vertex with the smallest tentative distance from start

# Dijkstra's example




	Distance	Parent
0	0	--
1	7	0
2	9	0
3	11	2
4	21	2
5	20	3

# Analysis of Dijkstra's algorithm

- How to implement?
  - Best path/parent array?
    - Runtime?
  - PQ?
    - Turns out to be very similar to Eager Prims
      - Storing paths instead of edges
    - Runtime?

# Dijkstra's example with negative edge weights



	Distance	Parent
0	0	--
1	10	0
2	5	0

**Incorrect!**


# Analysis of Dijkstra's algorithm

Dijkstra's is correct only when all edge weights  $\geq 0$

# Bellman-Ford's algorithm

- Set a distance value of `Double.POSITIVE_INFINITY` for all vertices
- Initialize a FIFO Q
- `distance[start] = 0`
- add start to Q
- While Q is not empty:
  - `cur = pop` a vertex from Q
  - For each non-parent neighbor x of cur:
    - Compute distance from start to x through cur
      - `distance[cur] + weight of edge between cur and x`
    - if computed distance < `distance[x]`
      - Update `distance[x]`
      - add x to Q if not already there

# Bellman-Ford's example with negative edge weights



	Distance	Parent
0	0	--
1	10	0
2	3	1

FIFO Q:

0  
1  
2

**Correct!**


# Analysis of Bellman-Ford's algorithm

Bellman-Ford's is correct even when there are negative edge weights in the graph but what about negative cycles?

- a negative cycle is a cycle with a negative total weight



# Bellman-Ford's example with a negative cycle



# Bellman-Ford's algorithm

- Set a distance value of Double.POSITIVE\_INFINITY for all vertices
- Initialize a FIFO Q
- $\text{distance}[\text{start}] = 0$
- add start to Q
- While Q is not empty **and no negative cycle has been detected**:
  - cur = pop a vertex from Q
  - For each non-parent neighbor x of cur:
    - Compute distance from start to x through cur
      - $\text{distance}[\text{cur}] + \text{weight of edge between cur and x}$
    - if computed distance < distance[x]
      - Update distance[x]
      - add x to Q if not already there
  - check for a negative cycle in the current Spanning Tree every v edges