



University of
Pittsburgh

Algorithms and Data Structures 2

CS 1501

Spring 2022

Sherif Khattab

ksm73@pitt.edu

(Slides are adapted from Dr. Ramirez's and Dr. Farnan's CS1501 slides.)

Announcements

- Upcoming deadlines:
 - Lab 10 due on 4/8
 - Homework 11 due on 4/11
 - Assignment 3 and 4 due on 4/18
 - Used to be one assignment

Previous lecture ...

- Network Bottleneck Finding Problem
 - Ford-Fulkerson Max-Flow Framework
 - augmenting path
 - backwards edges
 - Edmonds-Karp algorithm
 - using BFS to find augmenting paths

CourseMIRROR Reflections (most confusing)

- why is the max of the back edge the numerator of the forward edge?
- How/why we are able to take the backwards paths whenever the graph was originally directed
- Can any path from s to t be an augmenting path?
- How are bottlenecks defined? Are they flow edges that saturate early?
- Dijkstra algorithm but I will go over it by using past lectures
- Flow was mostly clear but doing another example would be helpful
- I was confused about how the runtime is calculated and what role maximum flow plays in it
- Going over backwards paths one more time and showing how they reroute the existing data flow would be helpful
- Edmonds Karp was most confusing.
- the ford example with 2000 iterations was confusing

CourseMIRROR Reflections (most interesting)

- How the max flow value is directly involved in the runtime of Ford-Fulkerson
- djikstras algorithm for shortest path explanation
- The data flow problem
- Residual graph
- The concept of using back edges
- Optimizing max flow through use of back edges
- adding backwards edges
- How backwards edges can help correct mistakes
- I found it interesting how BFS can help find any augmenting path
- The bottleneck problem in general was most interesting.
- ford fulkerson

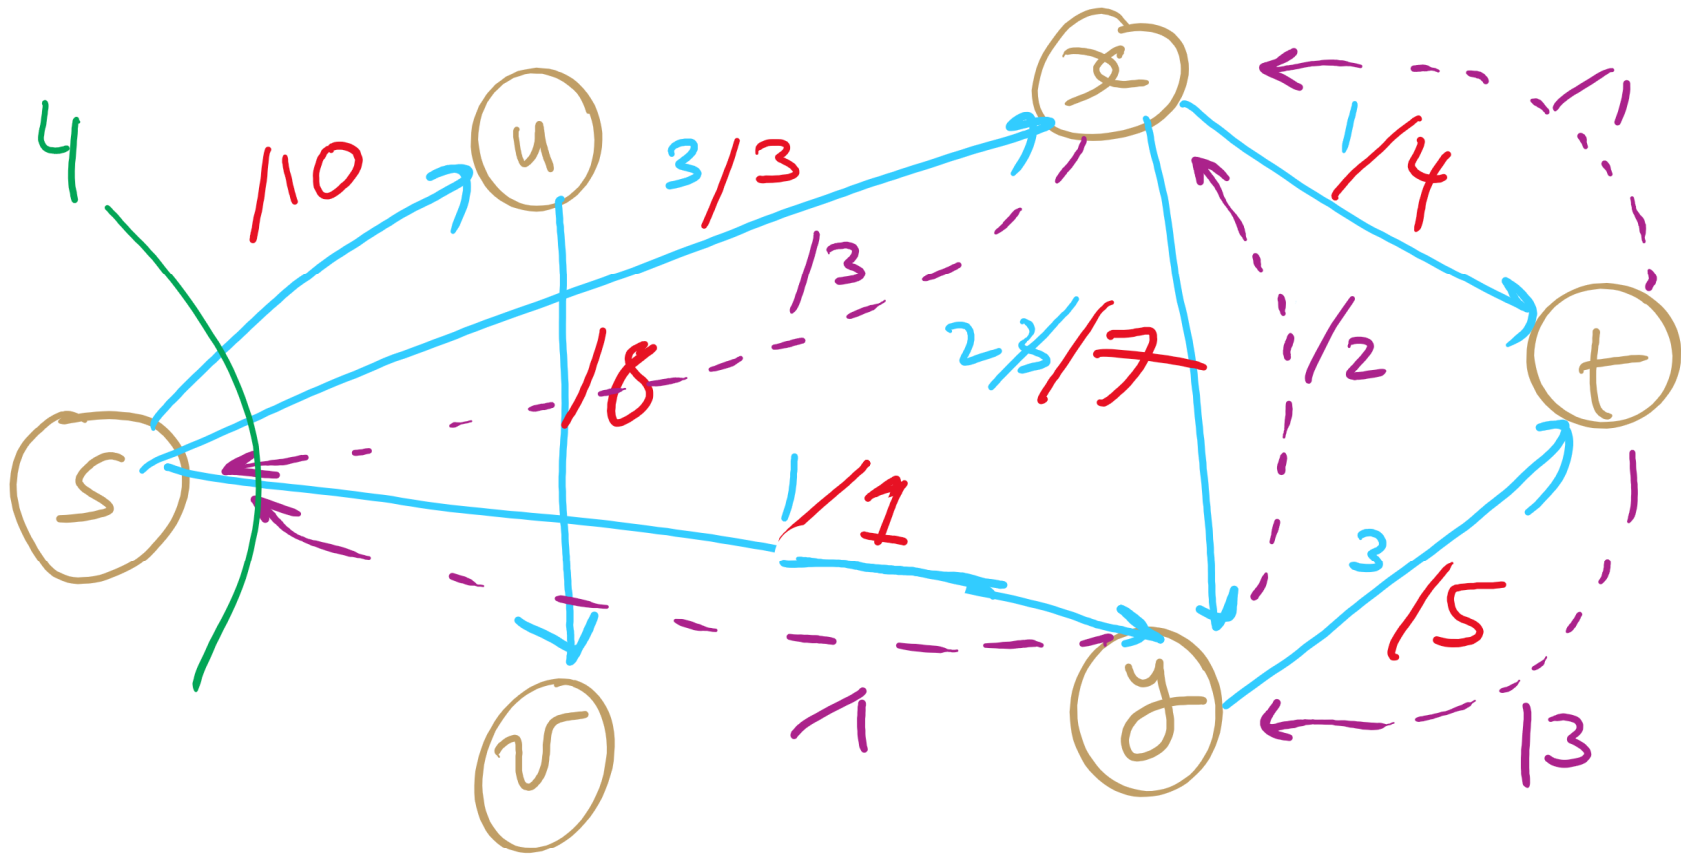
Problem of the Day: Finding Network Bottlenecks

- Let's assume that we want to send a large file from point A to point B over a computer network as fast as possible over multiple network links if needed
- Input:
 - A computer network
 - Network nodes and links
 - Links are labeled by link capacity in Mbps
 - Starting node and destination node
- Output:
 - The maximum network speed possible for sending a file from source to destination

But our flow graph is weighted...

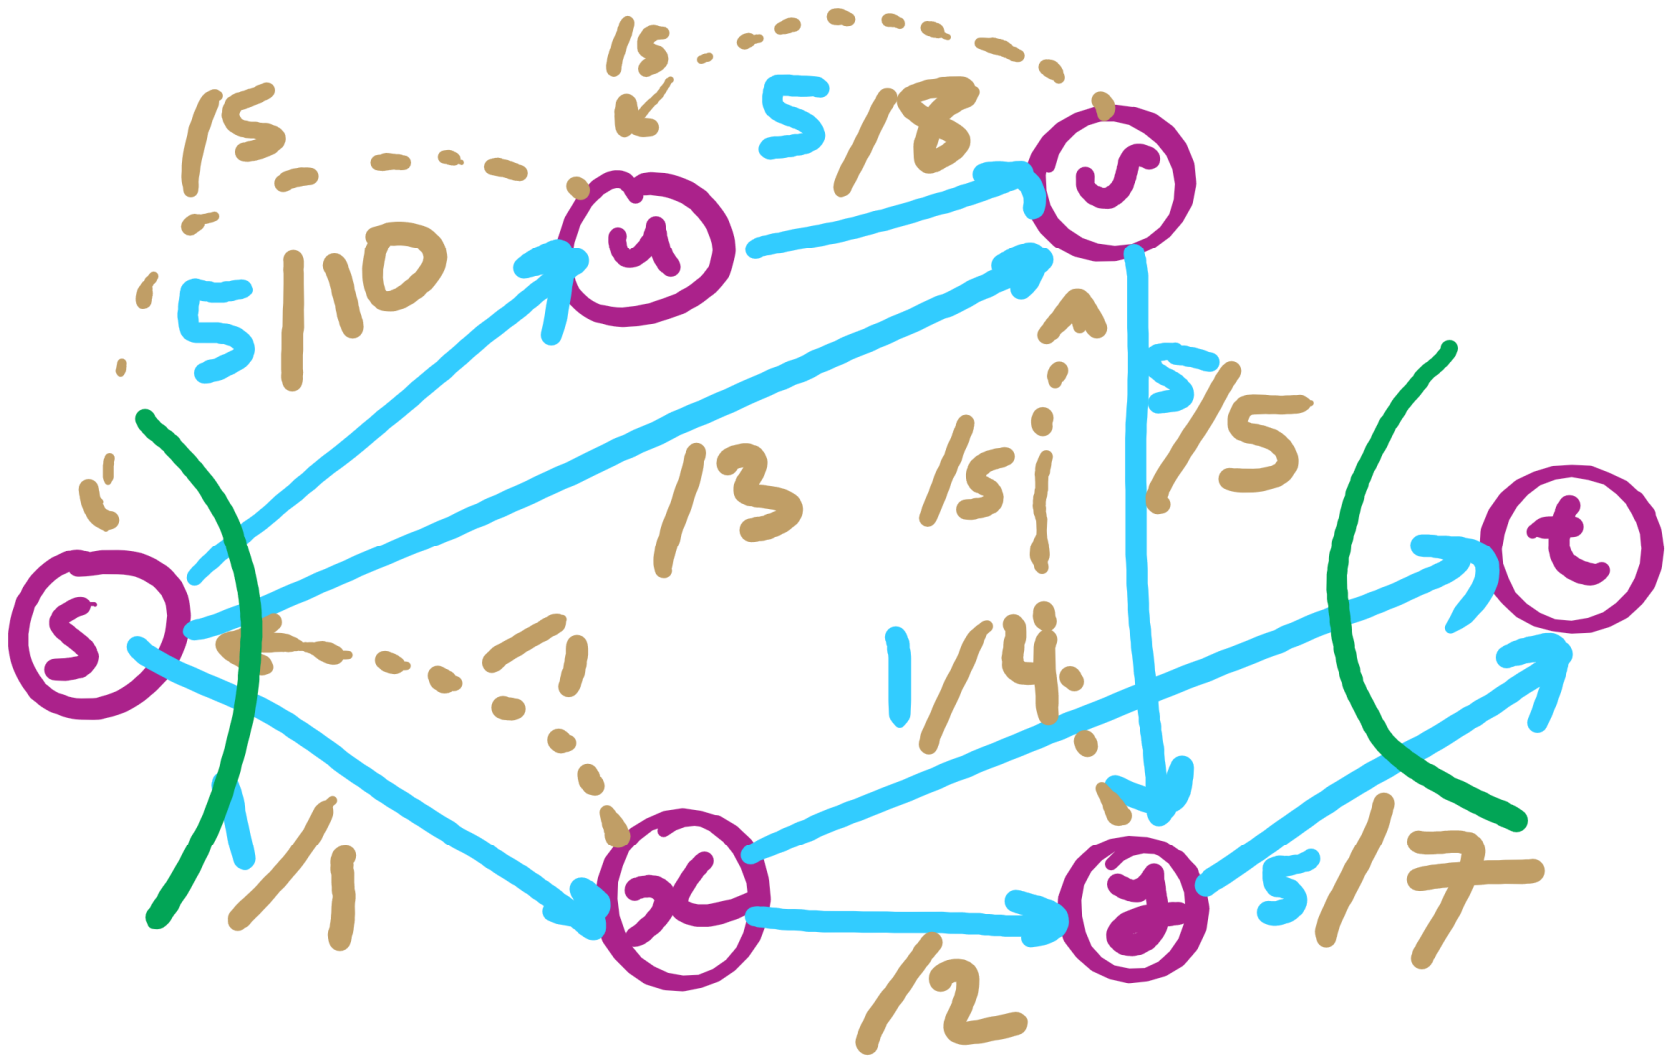
- Edmonds-Karp only uses BFS
 - Used to find spanning trees and shortest paths for *unweighted* graphs
 - Why do we not use some measure of priority to find augmenting paths?

PFS Example 1



$s \rightarrow x \rightarrow y \rightarrow t$
 $s \rightarrow y \rightarrow x \rightarrow t$

PFS Example 2



Flow edge implementation

- For each edge, we need to store:
 - Start point, the from vertex
 - End point, the to vertex
 - Capacity
 - Flow
 - Residual capacities
 - For forwards and backwards edges

FlowEdge.java

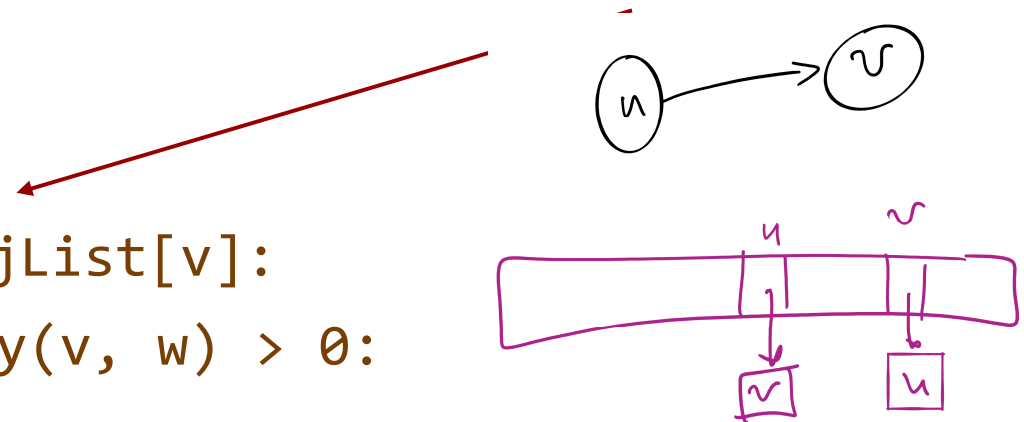
```
public class FlowEdge {  
    private final int v;           // from  
    private final int w;           // to  
    private final double capacity; // capacity  
    private double flow;           // flow  
  
    ...  
    public double residualCapacityTo(int vertex) {  
        if (vertex == v) return flow;  
        else if (vertex == w) return capacity - flow;  
        else throw new  
            IllegalArgumentException("Illegal endpoint");  
    }  
    ...  
}
```

BFS search for an augmenting path (pseudocode)

```
edgeTo = [|V|]
marked = [|V|]
Queue q
q.enqueue(s)
marked[s] = true
while !q.isEmpty():
    v = q.dequeue()
    for each (v, w) in AdjList[v]:
        if residualCapacity(v, w) > 0:
            if !marked[w]:
                edgeTo[w] = v;
                marked[w] = true;
                q.enqueue(w);
```

Each FlowEdge object is stored
in the adjacency list twice:

Once for its forward edge
Once for its backward edge



Value of maxflow

- Add up the flow increments in each iteration of Ford-Fulkerson
- Add up the edge **flows** out of source
- Add up the edge **flows** of the out of source

Follow-up Problem

- So, now we found the bottleneck *value*, but which edges define the found bottleneck?
 - *Why would you want to know those bottleneck edges?*

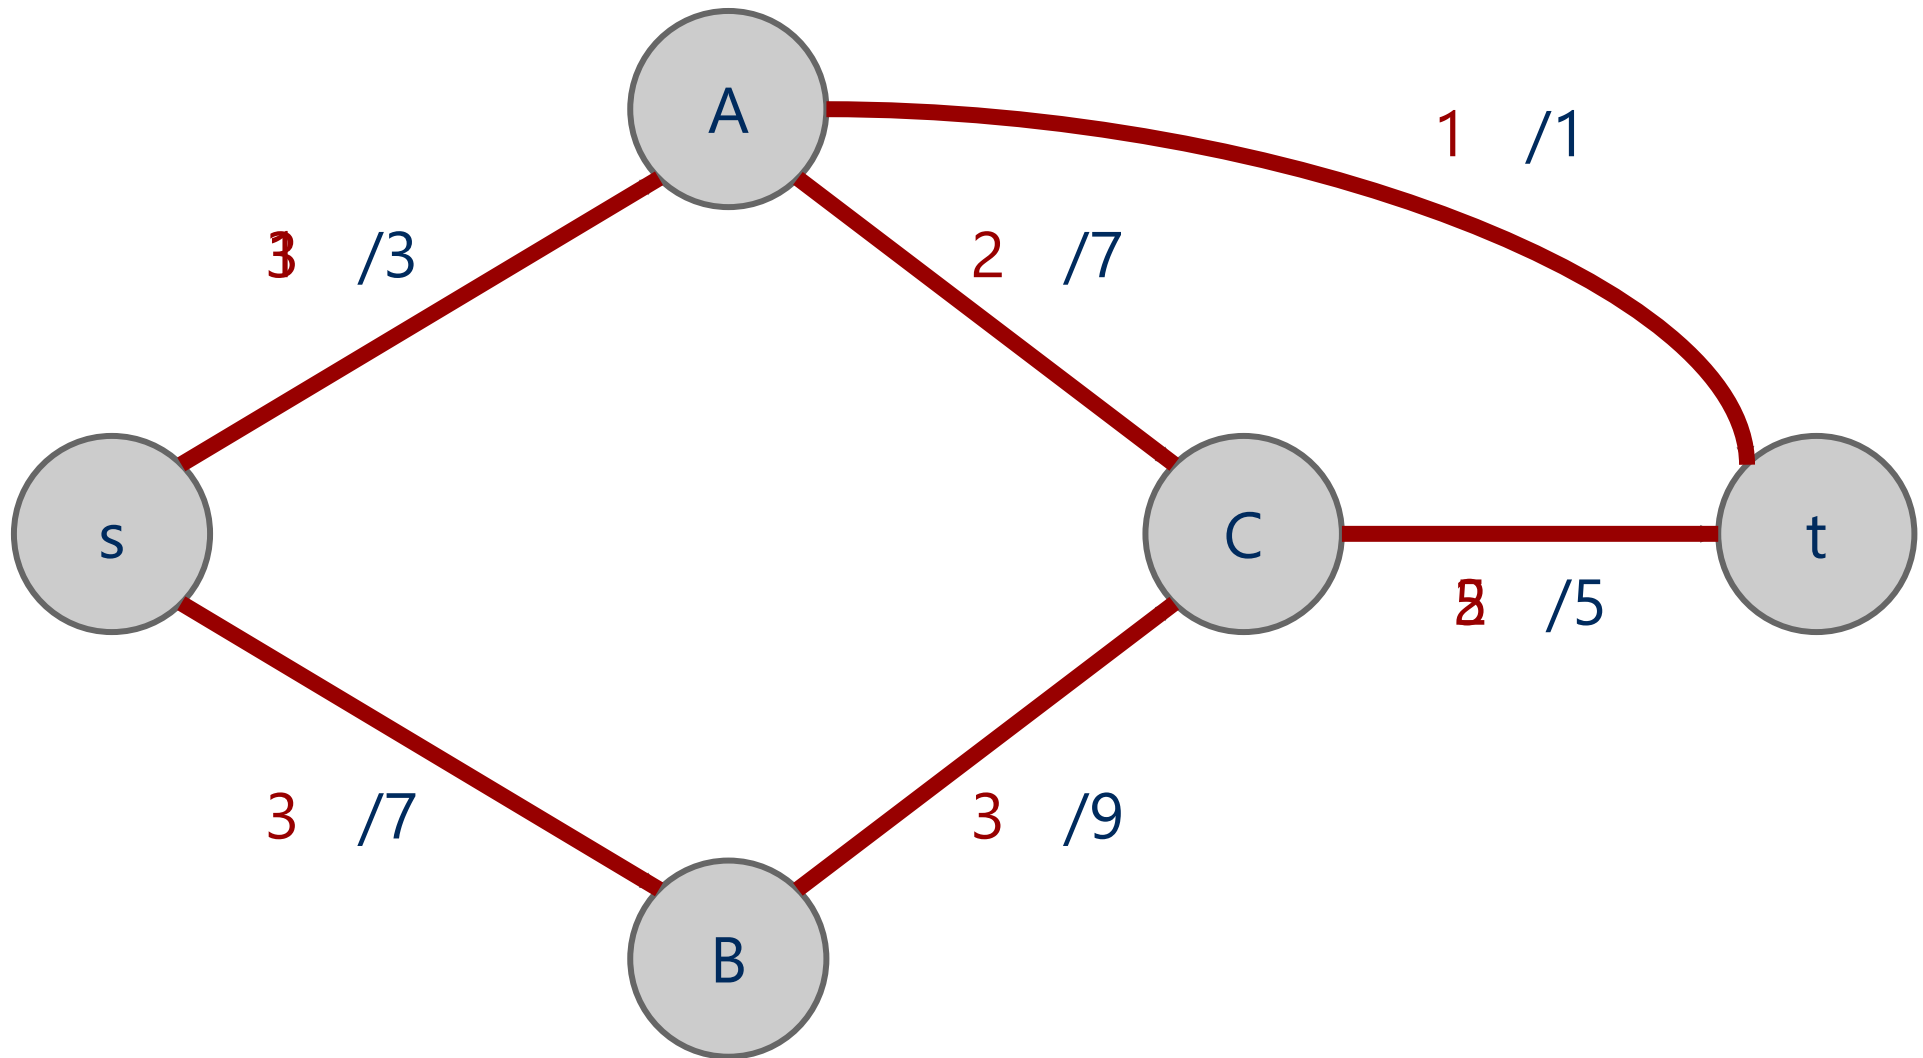
Let's separate the graph

- An st-cut on G is a set of edges in G that, if removed, will partition the vertices of G into two disjoint sets
 - One contains s
 - One contains t
- May be many st-cuts for a given graph
- Let's focus on finding the minimum st-cut
 - The st-cut with the smallest capacity
 - May not be unique

How do we find the min st-cut?

- We could examine residual graphs
 - Specifically, try and allocate flow in the graph until we get to a residual graph with no existing augmenting paths
 - A set of saturated edges will make a minimum st-cut

Min cut example



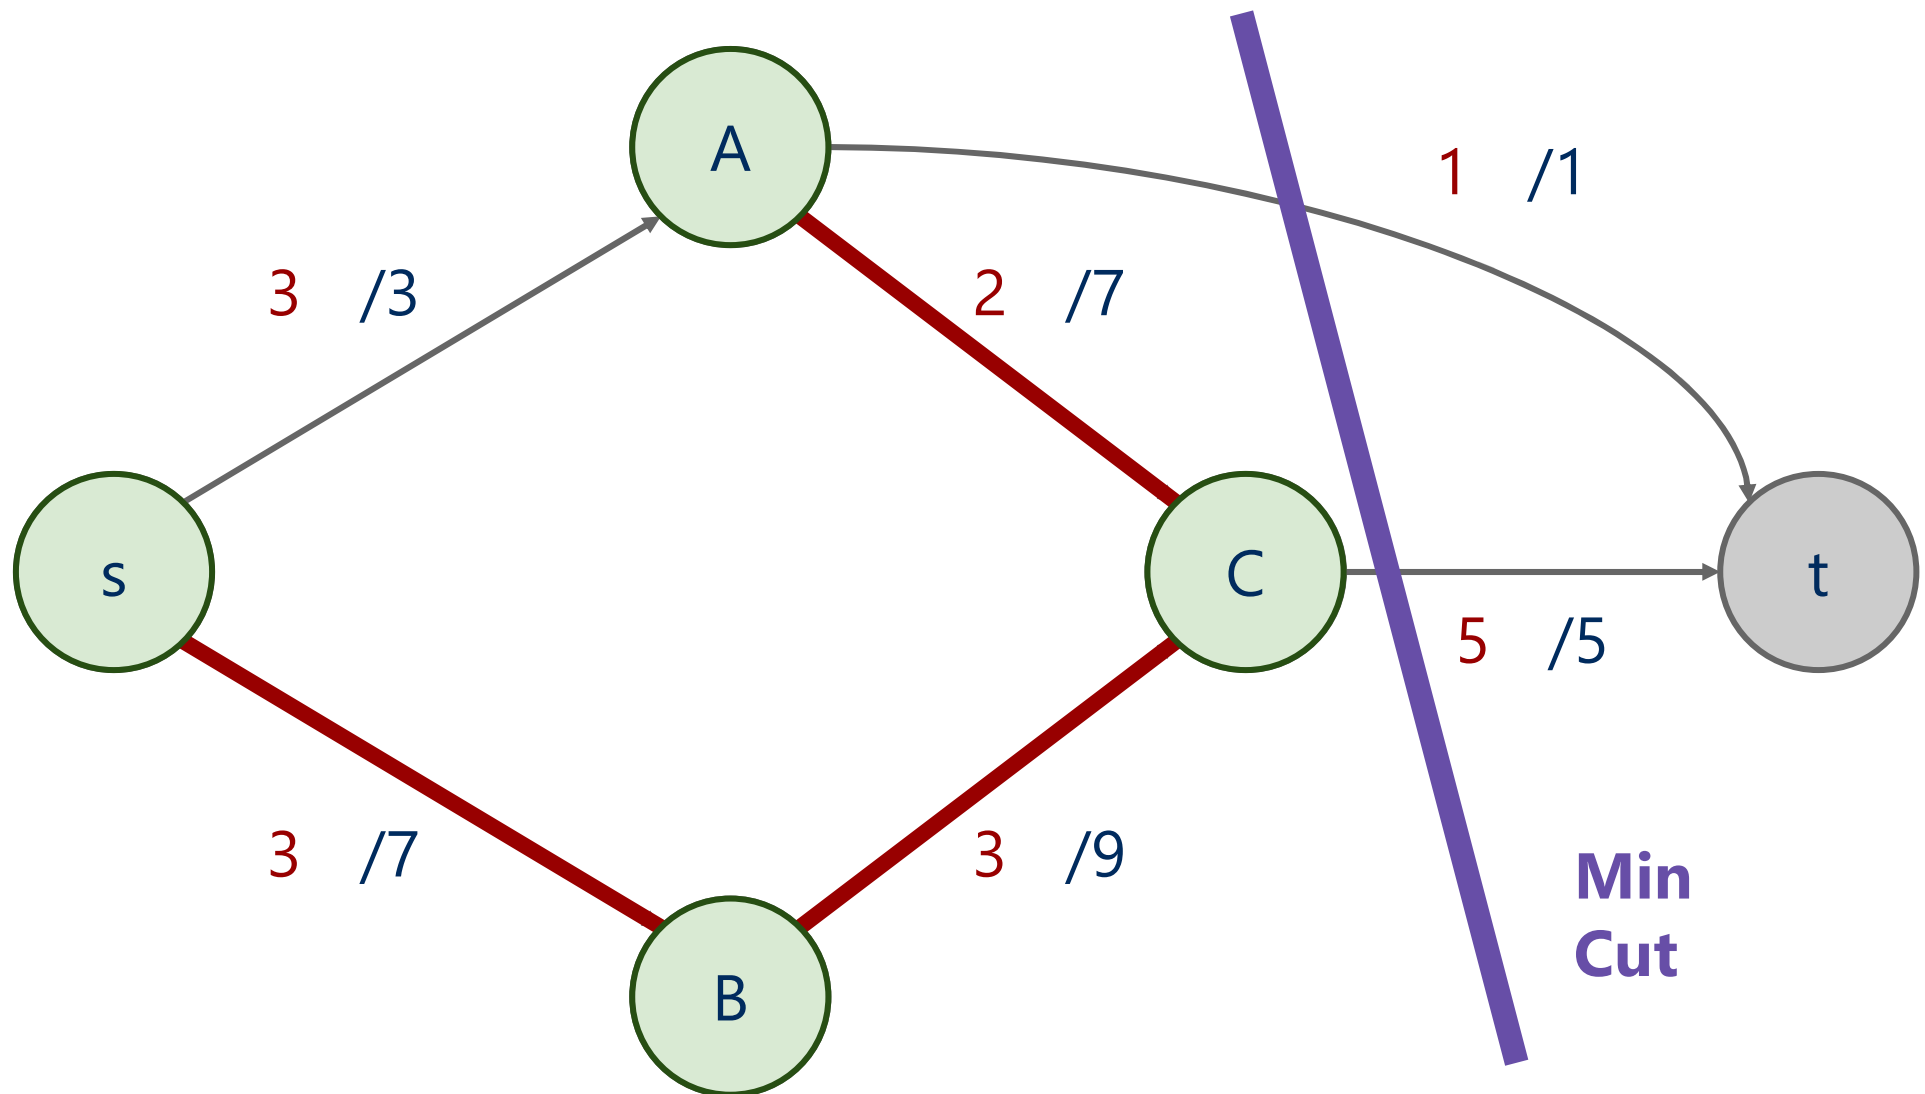
Max flow == min cut

- A special case of duality
 - I.e., you can look at an optimization problem from two angles
 - In this case to find the maximum flow or minimum cut
 - In general, dual problems do not have to have equal solutions
 - The differences in solutions to the two ways of looking at the problem is referred to as the *duality gap*
 - If the duality gap = 0, strong duality holds
 - Max flow/min cut uphold strong duality
 - If the duality gap > 0, weak duality holds

Determining a minimum st-cut

- First, run Ford Fulkerson to produce a residual graph with no further augmenting paths
- The last attempt to find an augmenting path will visit every vertex reachable from s
 - Edges with only one endpoint in this set comprise a minimum st-cut

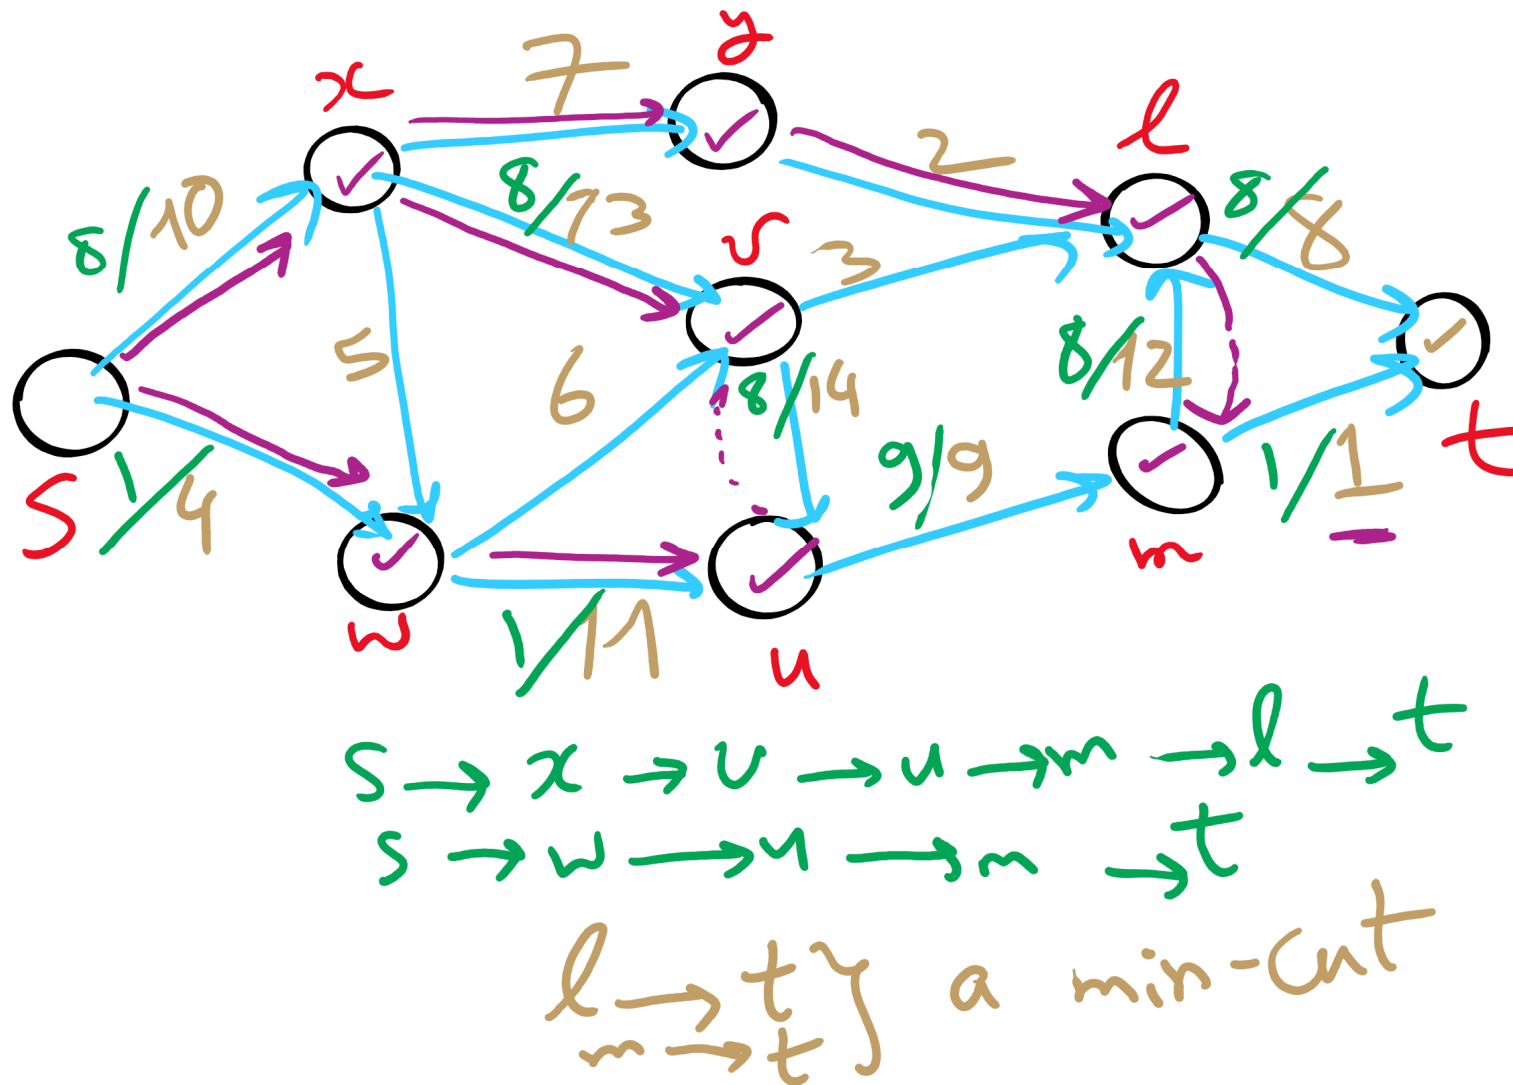
Determining the min cut



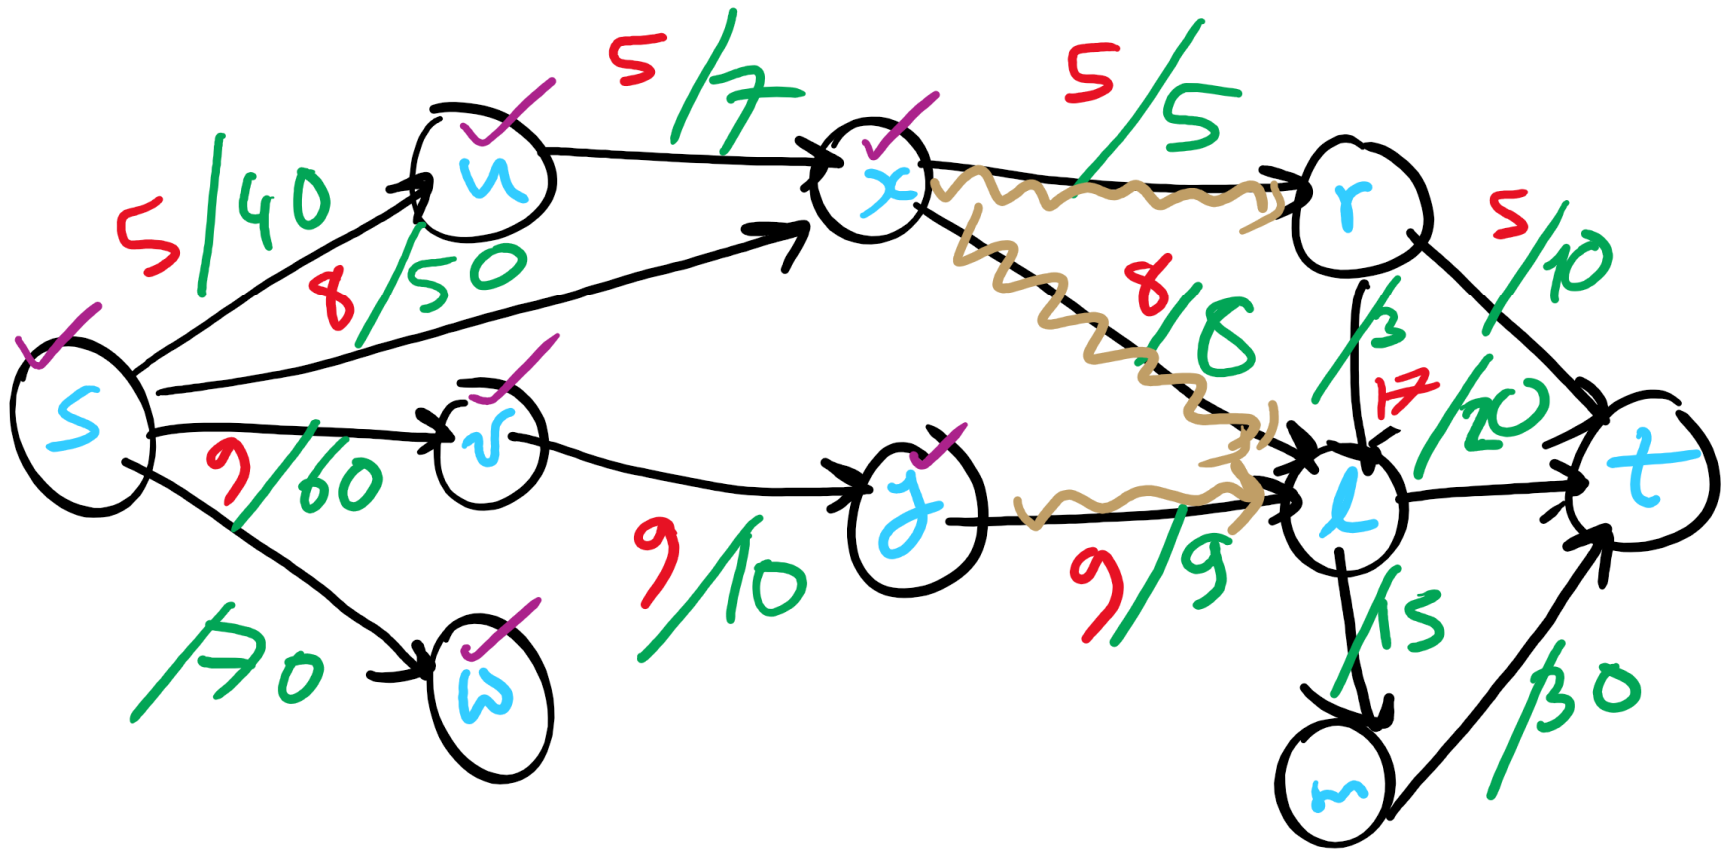
Max flow / min cut on unweighted graphs

- Is it possible?
- How would we measure the Max flow / min cut?
- What would an algorithm to solve this problem look like?

Min st-cut Example 1

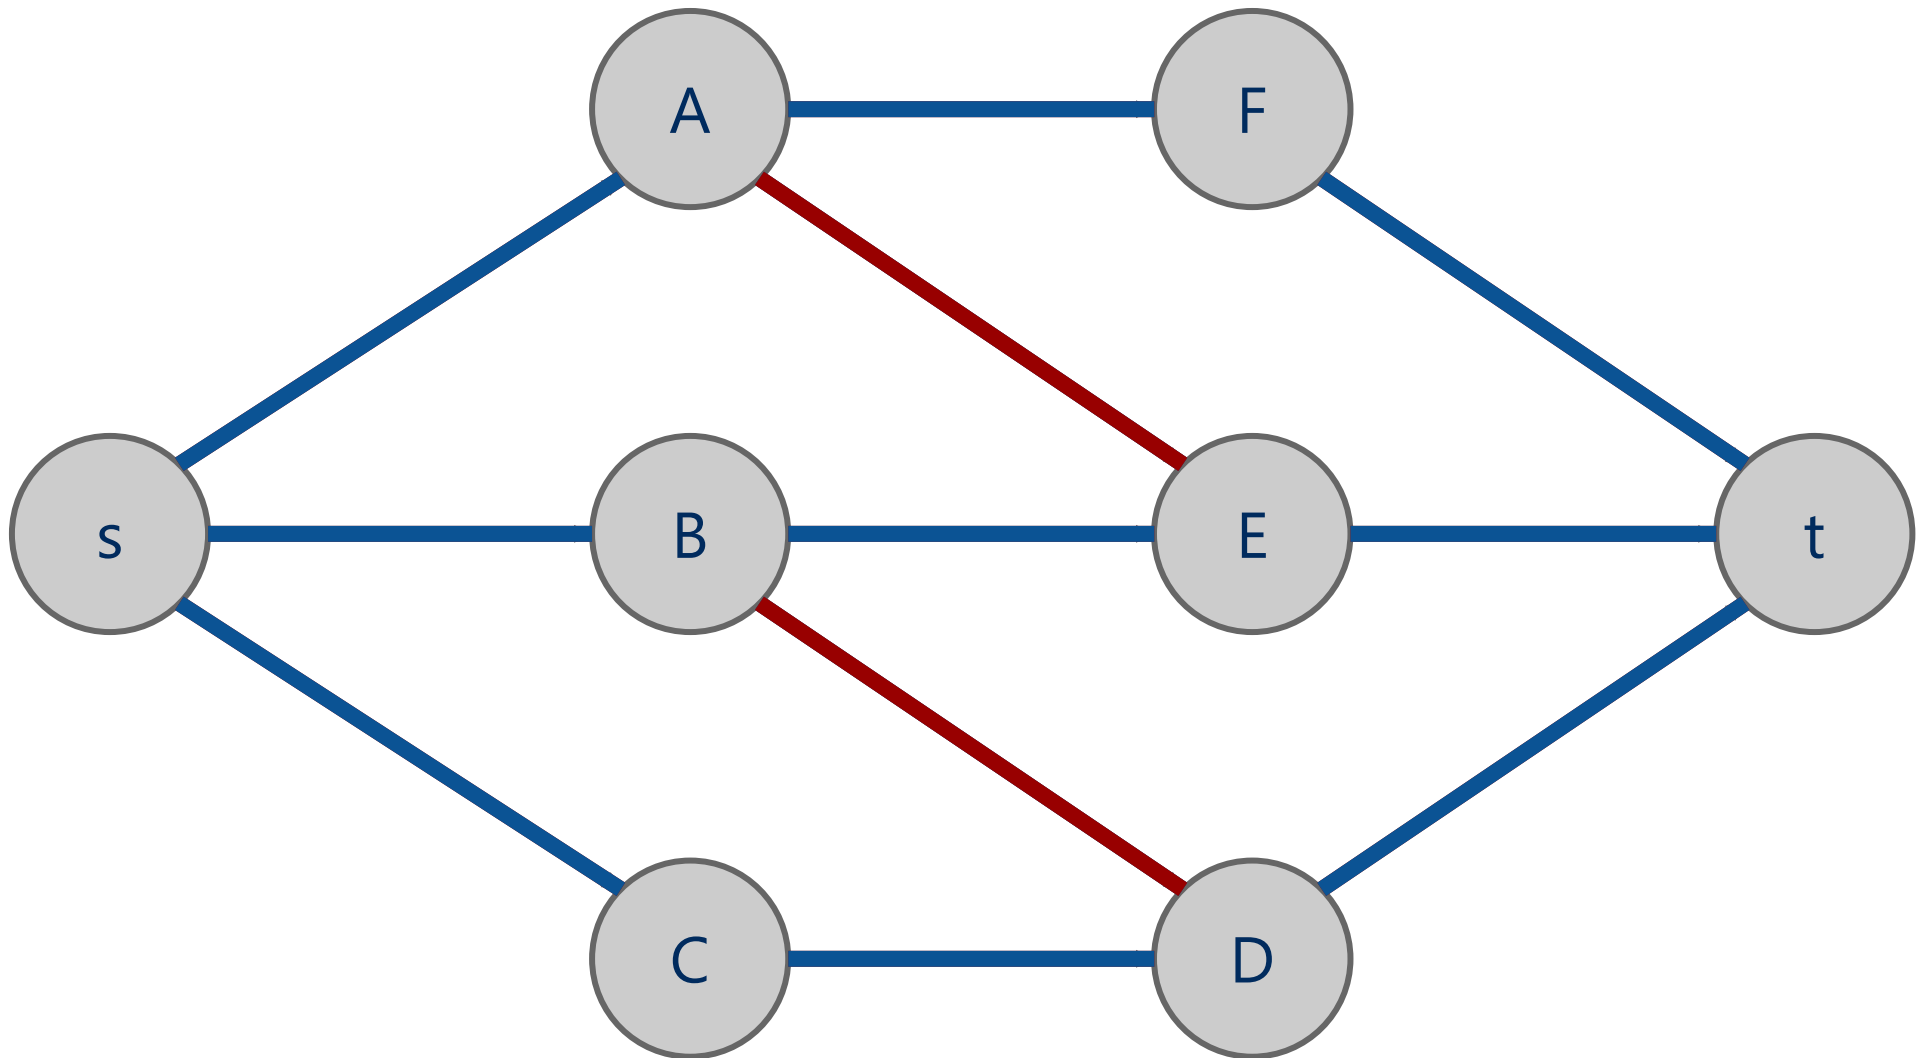


Min st-cut Example 2



$s \rightarrow u \rightarrow x \rightarrow r \rightarrow t$
 $s \rightarrow v \rightarrow y \rightarrow z \rightarrow t$
 $s \rightarrow x \rightarrow l \rightarrow t$

Unweighted network flow



Graph Algorithms Runtime

BFS	$\Theta(e+v)$	$\Theta(v^2)$
DFS	Adj. Lists	Adj Matrix
Connected Components		
Articulation Points		
Shortest Pathes based on # hops		
Prim's MST	$\Theta(v^3)$ naive	$\Theta(v^2)$ Best edge
		$\Theta(e \log e)$ PQ (Lazy)
		$\Theta(e \log v)$ Indexable PQ (Fiber)
Kruskal's MST	$\Theta(e(e+v))$ BFS/DFS for cycle detection	$\Theta(e \log e)$ UF weighted Trees
	$\Theta(v^2)$ Distance array	$\Theta(e \log v)$ Indexable PQ
Dijkstra's Shortest Pathes		
Ford-Fulkerson	$\Theta(f \cdot (e+v))$	
Edmonds-Karp (Max Flow & Min Cut)	$\Theta(e^2 \cdot v)$	

Problem of the Day: Integer multiplication

- Input: two n -bit integers x and y
 - n can be really large!
 - e.g., 2048 bits
 - Why do we need such big numbers?
 - The input size is $\theta(n)$
 - The values of x and y are exponential in the input size!
- Output:
 - $x * y$

Yeah, but the processor has a MUL instruction

- Assuming x86
- Given two 32-bit integers, MUL will produce a 64-bit integer in a few cycles
- What about when we need to multiply large ints?
 - VERY large ints?
 - RSA keys should be 2048 bits
 - Back to grade school...

Gradeschool algorithm on binary numbers

```

              10100000100
x             11000101111
-----
              10100000100
              101000001000
              1010000010000
              10100000100000
              000000000000000
              1010000010000000
              0000000000000000
              00000000000000000
              000000000000000000
              1010000010000000000
              10100000100000000000
-----
            111110000001110111100
```

OK, I'm guessing we all knew that...

- What is the runtime of this multiplication algorithm?
 - For 2 n-digit numbers:
 - n^2

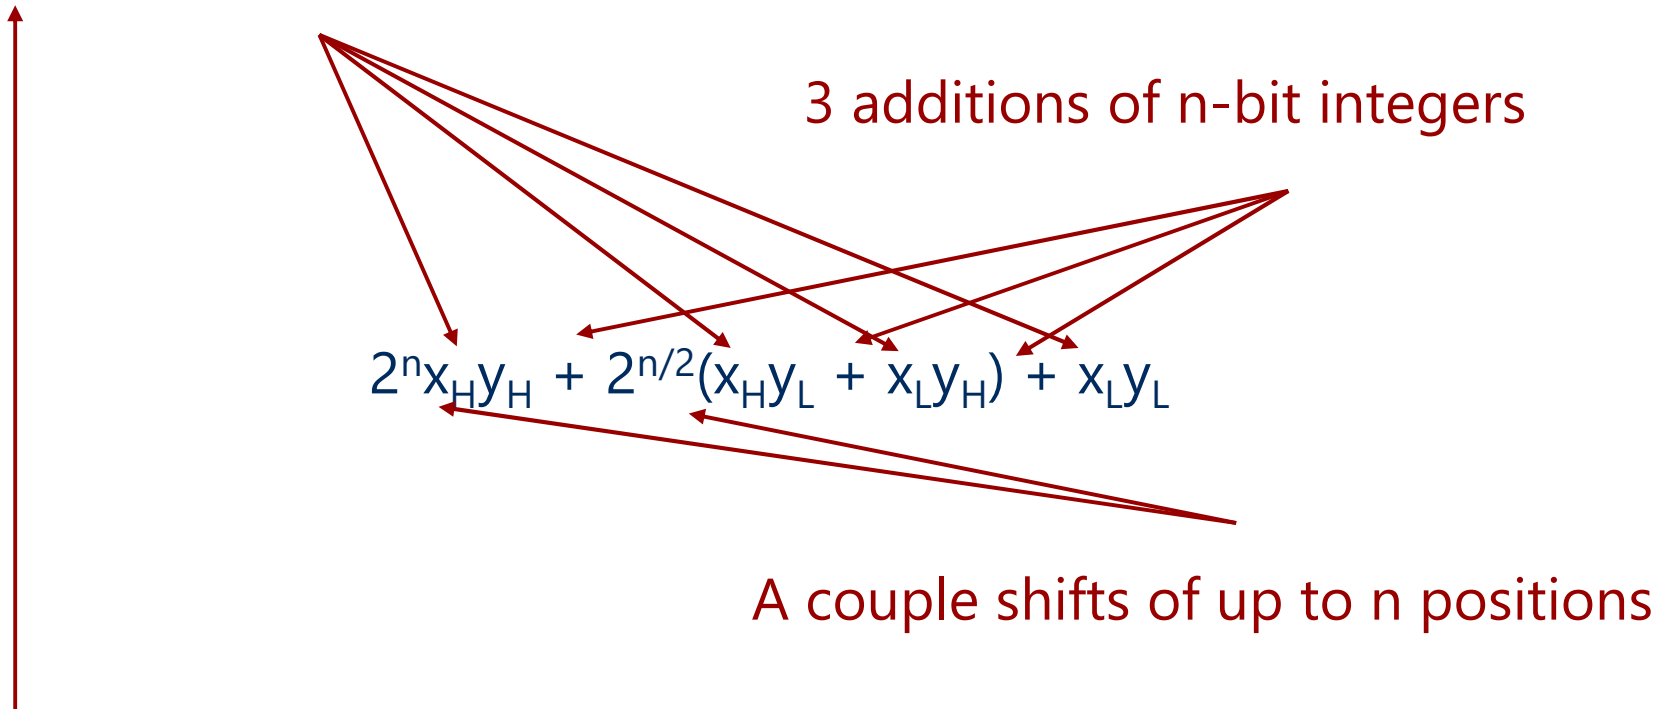
How can we improve our runtime?

- Let's try to divide and conquer:
 - Break our n -bit integers in half:
 - $x = 1001011011001000$, $n = 16$
 - Let the high-order bits be $x_H = 10010110$
 - Let the low-order bits be $x_L = 11001000$
 - $x = 2^{n/2}x_H + x_L$
 - Do the same for y
 - $x * y = (2^{n/2}x_H + x_L) * (2^{n/2}y_H + y_L)$
 - $x * y = 2^n x_H y_H + 2^{n/2}(x_H y_L + x_L y_H) + x_L y_L$

So what does this mean?

Multiplying
2 4-bit \Rightarrow 8-bit
2 $n/2$ -bit \Rightarrow n -bit
2 n -bit $=$ $2n$ -bit

4 multiplications of $n/2$ bit integers



Actually 16 multiplications of $n/4$ bit integers (plus additions/shifts)

Actually 64 multiplications of $n/8$ bit integers (plus additions/shifts)

...

Karatsuba's algorithm

- By reducing the number of recursive calls (subproblems), we can improve the runtime

- $x * y = 2^n x_H y_H + 2^{n/2} (x_H y_L + x_L y_H) + x_L y_L$

M1

M2

M3

M4

- We don't actually need to do both M2 and M3
 - We just need the sum of M2 and M3
 - If we can find this sum using only 1 multiplication, we decrease the number of recursive calls and hence improve our runtime

Karatsuba craziness

- $M1 = x_h y_h$; $M2 = x_h y_l$; $M3 = x_l y_h$; $M4 = x_l y_l$;
- The sum of all of them can be expressed as a single mult:
 - $M1 + M2 + M3 + M4$
 - $= x_h y_h + x_h y_l + x_l y_h + x_l y_l$
 - $= (x_h + x_l) * (y_h + y_l)$
- Lets call this single multiplication M5:
 - $M5 = (x_h + x_l) * (y_h + y_l) = M1 + M2 + M3 + M4$
- Hence, $M5 - M1 - M4 = M2 + M3$
- So: $x * y = 2^n M1 + 2^{n/2} (M5 - M1 - M4) + M4$
 - Only 3 multiplications required!
 - At the cost of 2 more additions, and 2 subtractions

Karatsuba Example

$$\begin{array}{l}
 x = \underbrace{977540}_{x_H} \underbrace{1238}_{x_L} \\
 y = \underbrace{145970}_{y_H} \underbrace{1003}_{y_L}
 \end{array}$$

$n = 10$

$$M_1 = x_H \cdot y_H$$

$$M_2 = x_L \cdot y_L$$

$$M_3 = (x_L + x_H)(y_L + y_H)$$

$$\begin{aligned}
 x * y = & M_1 \text{ shifted left by } n \text{ positions} \\
 & + (M_3 - M_1 - M_2) \text{ shifted left by } \frac{n}{2} \text{ positions} \\
 & + M_2
 \end{aligned}$$

Large integer multiplication in practice

- Can use a hybrid algorithm of grade school for large operands, Karatsuba's algorithm for VERY large operands
 - Why are we still bothering with grade school at all?

Is this the best we can do?

- The Schönhage–Strassen algorithm
 - Uses Fast Fourier transforms to achieve better asymptotic runtime
 - $O(n \log n \log \log n)$
 - Fastest asymptotic runtime known from 1971-2007
 - Required n to be astronomical to achieve practical improvements to runtime
 - Numbers beyond $2^{2^{15}}$ to $2^{2^{17}}$
- Fürer was able to achieve even better asymptotic runtime in 2007
 - $n \log n 2^{O(\log^* n)}$
 - No practical difference for realistic values of n

$\log^* n$ = # times we can compute $\log n$ until the result ≤ 1

Please submit your reflections by using the CourseMIRROR App

If you are having a problem with CourseMIRROR, please send an email to coursemirror.development@gmail.com

8/29/2022