



University of
Pittsburgh

Algorithms and Data Structures 2

CS 1501



Spring 2023

Sherif Khattab

ksm73@pitt.edu

(Slides are adapted from Dr. Ramirez's and Dr. Farnan's CS1501 slides.)

Announcements

- Midterm grades posted (out of 60)
 - Question reattempts to get up to **7 points** back
 - Please use GradeScope's Regrade Requests for each question individually **due on Monday 4/17 at 11:59 pm**
- Upcoming Deadlines
 - Lab 10: Tuesday 4/11 @ 11:59 pm
 - Homework 11: **next Friday** @ 11:59 pm
 - Assignment 4: Friday 4/14 @ 11:59 pm
 - Support video and slides on Canvas + Solutions for Labs 8 and 9

Previous lecture

- Minimum Spanning Tree (MST)
 - Prim's MST algorithm
 - naïve implementation
 - Best Edges array implementation
 - using a min-heap
 - Kruskal's MST algorithm

This Lecture

- Weighted Shortest Paths problem
 - Dijkstra's single-source shortest paths algorithm
 - Bellman-Ford's shortest paths algorithm

Kruskal's MST algorithm

- Insert all e edges into a PQ
- T = an empty set of edges
- Repeat until T contains **$v-1$** edges
 - Remove a min edge from the PQ
 - Add the edge to T if the edge does not create a cycle in T
- return T


Runtime of Kruskal's MST algorithm

- Instead of building up the MST starting from a single vertex, we build it up using edges all over the graph
- How do we efficiently implement cycle detection?
 - BFS/DFS
 - $v + e$
 - Union/Find data structure (not covered)
 - $\log v$

Kruskal's Runtime

- e iterations
 - $\text{removeMin} \rightarrow \log e$
 - Cycle detection
 - $v + e$ using DFS/BFS
 - $\log v$ using Union/Find
- Total runtime = $\Theta(e \log e)$
- Assuming connected graph
 - $v - 1 \leq e \leq v^2$
 - $\log v \leq \log e \leq 2 \log v$
 - $\log e = \Theta(\log v)$
- Total runtime = $\Theta(e \log e) = \Theta(e \log v)$
- Same runtime as Prim's

Problem of the Day: Weighted Shortest Paths



1.6M vertices, 3.8M arcs, travel time metric.

Source: <https://www.cs.princeton.edu/courses/archive/spr09/cos423/Lectures/reach-mit.pdf>₈

Problem of the Day: Weighted Shortest Paths

- **Input:** starting and destination addresses and a road network
 - Road segments and intersections
 - Road segments are labeled by **travel time**
- **Output:**
 - A **shortest path** from starting address to destination address


Dijkstra's algorithm: Data Structures and Initialization

- `distance[]`: **best known** shortest distance from start to each vertex
- `distance[start] = 0`
- `distance[x] = Double.POSITIVE_INFINITY` for other vertices

Dijkstra's algorithm: Data Structures and Initialization

- $cur = start$
- While destination not visited:
 - For each **unvisited** neighbor x of cur
 - Compute **shortest** distance from start to x **through cur**
 - $= distance[cur] + \text{weight of edge from cur to } x$
 - Update $distance[x]$ if distance through $cur < distance[x]$
 - Mark cur as visited
 - $cur =$ an unvisited vertex with the smallest distance

Dijkstra's example



	Distance	Parent
0	0	--
1	7	0
2	9	0
3	11	2
4	21	2
5	20	3

Notes


- The distance array keeps track of the **best path** from start
 - Compare to **best edge** array in Prim's
- Once a vertex is **visited**, its distance value **doesn't change**
- **Parent** array used to construct a shortest path

Analysis of Dijkstra's algorithm

- Depends on implementation!
 - Distance and parent arrays? **Theta(v^2)**
 - PQ?
 - very similar to **Eager Prim's** **Theta($e \log v$)**
 - Storing best paths instead of best edges
- This is worst-case runtime
- Algorithm may **stop earlier** when destination visited
- Order of selecting vertices matters!

Dijkstra's Real-World Optimizations

- **Real-world road networks:**
 - millions of vertices and edges
 - want fast response (< 100 ms)




1.6M vertices, 3.8M arcs, travel time metric.

Source: <https://www.cs.princeton.edu/courses/archive/spr09/cos423/Lectures/reach-mit.pdf>

Dijkstra's is too slow for such big graphs


- Dijkstra's will visit so many **not needed** vertices



Searched area

Optimization 1: Bidirectional Search

- start two instances of Dijkstra's possibly **in parallel**
 - from source on **original** graph
 - from destination, on **reverse** graph
- When processing an edge to a vertex visited by the other instance, update **shortest known** distance between start and destination
- **Stop** when tops of both heaps give a distance \geq shortest known




forward search / reverse search

Optimization 2: A* Search

- Use **lower-bound estimates** for the distance of the **rest** of the path to destination
- Modified Dijkstra's
 - Pick vertex with minimum **distance[v] + estimate[v]**

- Lower-bound estimates using **landmarks** and **triangular inequality**



- Requires **preprocessing** to compute and store distances from each vertex to each landmark


Negative Edge Weights

Modeling of some problems results in graphs with **negative** edge weights

Example: Planning with **uncertainty**

- Find a path with **highest probability** of reaching a goal from a starting state
- **vertices**: milestones; **edges**: actions; **edge weights**: probabilities
- Find a path with **highest product** of edge weights
- How to model that as a shortest path problem?
 - log of product is sum of logs
 - maximize x means minimize -x
 - replace each edge weight p by $(-1 * \log p)$
 - find a shortest path in the resulting graph!

Dijkstra's example with negative edge weights



	Distance	Parent
0	0	--
1	10	0
2	5	0

Incorrect!

Dijkstra's algorithm is incorrect with negative edge weights

Dijkstra's is correct only when all edge weights ≥ 0

Bellman-Ford's algorithm: Data Structures and Initialization

- $\text{distance}[v] = \text{Double.POSITIVE_INFINITY}$
 - for all vertices except start
- $\text{distance}[\text{start}] = 0$

Bellman-Ford's algorithm

- Repeat $v-1$ times
 - For each vertex cur :
 - For each neighbor x of cur :
 - Compute shortest distance from start to x via cur
 - $= \text{distance}[cur] + \text{weight of } (cur, x)$
 - if computed distance $< \text{distance}[x]$
 - Update $\text{distance}[x]$ and $\text{parent}[x]$


Runtime of Bellman-Ford's

- Repeat $v-1$ times
 - For each vertex cur :
 - For each neighbor x of cur :
 - Compute shortest distance from start to x via cur
 - $= \text{distance}[cur] + \text{weight of } (cur, x)$
 - if computed distance $< \text{distance}[x]$
 - Update $\text{distance}[x]$ and $\text{parent}[x]$
- Runtime?
 - $O(v \cdot e)$

Bellman-Ford's algorithm: an optimization

- Initialize a FIFO Q; add start to Q
- While Q is not empty:
 - $cur = \text{pop a vertex from } Q$
 - For each neighbor x of cur :
 - Compute shortest distance from start to x via cur
 - $= \text{distance}[cur] + \text{weight of } (cur, x)$
 - if computed distance $< \text{distance}[x]$
 - Update $\text{distance}[x]$ and $\text{parent}[x]$
 - add x to Q if not already there

Bellman-Ford's example with negative edge weights



	Distance	Parent
0	0	--
1	10	0
2	3	1

FIFO Q:

0
1
2

Correct!

Analysis of Bellman-Ford's algorithm


- Bellman-Ford's is correct even when there are negative edge weights in the graph but what about negative cycles?
 - a **negative cycle** is a cycle with a negative total weight

Negative cycles

- **Detecting** such cycle is needed:
 - Bellman-Ford **won't terminate** if a negative cycle exists
 - Some problems can be solved by detecting a negative cycle
- Example: Finding **arbitrage** in currency trade
 - **vertices**: currencies; **edge weights**: exchange rates;
 - **goal**: find a cycle with a **product** of exchange rates that is > 1
 - \log of product is sum of logs
 - maximize x means minimize $-x$
 - Replace each edge weight r by $(-1 * \log r)$
- a cycle with a product $> 1 \rightarrow$ a **negative cycle** in the resulting graph

Bellman-Ford's example with a negative cycle

- Can you find a negative cycle?



Find a negative cycle reachable from start

- Repeat $v-1$ times
 - For each vertex cur :
 - For each neighbor x of cur :
 - Compute shortest distance from start to x via cur
 - $= \text{distance}[cur] + \text{weight of } (cur, x)$
 - if computed distance $< \text{distance}[x]$
 - Update $\text{distance}[x]$ and $\text{parent}[x]$
- If **another iteration** results in update of $\text{distance}[v]$ for a vertex v , then v is in a negative cycle
 - can find the cycle using parent values starting from $\text{parent}[v]$

Find a negative cycle reachable from start

- May be able to detect a negative cycle **earlier**
- Build a graph using **parent to child links** set by Bellman-Ford's
- Modify **DFS** to detect if a cycle exists
 - if a **neighbor** already **visited** and is **on** the runtime **stack**
 - we have a **cycle**
 - follow parent links until back to current node
 - add up edge weights
 - if negative stop; otherwise continue